

The AWS Access Request Service (Phase-1) enables employees to request temporary AWS access using natural language. The system uses an LLM to reduce incomplete requests by suggesting services/actions and asking a single follow-up question when mandatory details (such as resource ARNs) are missing. All access requests follow a strict approval workflow and generate IAM policies and AWS CLI commands without automatic execution.

## 2.2 Core Components

### 2.2.1 Access Request API

- Entry point for all access requests
- Authenticated via existing SSO
- Accepts request reason and AWS account
- Orchestrates LLM calls and approval flow

## 2.2.2 LLM Assist Module

- Interprets request reason
- Suggests AWS services and action groups
- Determines if follow-up information is required
- Asks only one follow-up question at a time
- Produces deterministic JSON output

## 2.2.3 Access Request Store

- Persists access request data
- Maintains request lifecycle status
- Serves as the system's source of truth

## 2.2.4 Approval Workflow

- Manager approval for business validation
- DevOps approval for technical validation
- Manual execution of generated CLI commands

## 2.2.5 IAM Policy & CLI Generator

- Generates IAM policy JSON
- Generates AWS CLI command text
- Does not execute commands (Phase-1 constraint)

## 2.2.6 Audit Logging

- Captures all system events
- Immutable record for governance and compliance

---

## 3. Low Level Design (LLD)

### 3.1 Request Creation Flow

1. Employee submits an access request with a reason.
  2. Access Request API invokes LLM Assist Module.
  3. LLM suggests AWS services and action groups.
  4. If required information is missing:
    - LLM returns needFollowup = true
    - System asks a single follow-up question.
  5. User responds to follow-up.
  6. LLM re-evaluates and finalizes suggestions.
  7. Access request is stored with status CREATED.
  8. Audit log entry is recorded.
- 

### **3.2 Approval Flow**

#### **Manager Approval**

- Reviews request intent
- Approves or rejects with reason
- Status updated accordingly

#### **DevOps Approval**

- Reviews access scope and generated CLI
  - Approves or rejects
  - On approval, policy and CLI command are generated
- 

### **3.3 IAM Policy Generation Flow**

- Inputs:
    - AWS services
    - IAM actions
    - Resource ARNs
  - Outputs:
    - IAM policy JSON
    - AWS CLI command (text only)
-

## 4.

### Database Design (Phase-1)

#### 4.1 Entity Overview

Table Name	Purpose
users	Employee, Manager, DevOps identities
services	Allowed AWS services
action_groups	Logical permission groups
aws_actions	IAM actions per group
access_requests	Core request data
approvals	Manager & DevOps decisions
Audit_logs	Governance and audit trail
Suggestion	To give the follow up question
follow up	

---

### 5. DBMS Tables & SQL DDL

#### 5.1 user: -NOT REQ(INSTEAD OF USER ID USE EMAIL)

```
CREATE TABLE access_requests (
    id BIGINT PRIMARY KEY AUTO_INCREMENT,
    requester_email VARCHAR(255) NOT NULL,
    aws_account ENUM('ZINKA','DIVUM') NOT NULL,
    reason TEXT NOT NULL,
    services JSON NOT NULL,
    resource_arns JSON,
    status ENUM(
```

```
'CREATED',
'MANAGER_APPROVED',
'DEVOPS_APPROVED',
'ACCESS_GRANTED',
'REJECTED'
) NOT NULL,  
  
duration_hours INT DEFAULT 24,  
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
);
```

---

## 5.2 services

```
CREATE TABLE services (
    id BIGINT PRIMARY KEY AUTO_INCREMENT,
    service_name VARCHAR(100) NOT NULL
);
```

---

## 5.3 action\_groups

```
CREATE TABLE action_groups (
    id BIGINT PRIMARY KEY AUTO_INCREMENT,
    service_id BIGINT NOT NULL,
    Description VARCHAR(100) NOT NULL,
    group_name VARCHAR(100) NOT NULL,
    FOREIGN KEY (service_id) REFERENCES services(id)
);
```

---

## 5.4 aws\_actions

Changes made :  
Action\_groups

```
CREATE TABLE action_groups (
    id BIGINT PRIMARY KEY AUTO_INCREMENT,
    service_id BIGINT NOT NULL,
    group_name VARCHAR(100) NOT NULL
);
```

aws\_actions:

```
CREATE TABLE aws_actions (
    id BIGINT PRIMARY KEY AUTO_INCREMENT,
    aws_action VARCHAR(100) NOT NULL UNIQUE
);
```

junction table

```
CREATE TABLE action_group_actions (
    action_group_id BIGINT NOT NULL,
    aws_action_id BIGINT NOT NULL,
    PRIMARY KEY (action_group_id, aws_action_id),
    FOREIGN KEY (action_group_id) REFERENCES action_groups(id),
    FOREIGN KEY (aws_action_id) REFERENCES aws_actions(id)
);
```

---

## 5.5 access\_requests

```
CREATE TABLE access_requests (
    id BIGINT PRIMARY KEY AUTO_INCREMENT,
    requester_id BIGINT NOT NULL,
    aws_account ENUM('ZINKA','DIVUM') NOT NULL,
    reason TEXT NOT NULL,
    services JSON NOT NULL,
    resource_arns JSON,
    status ENUM(
        'CREATED',
```

```

'MANAGER_REJECTD',
'DEVOPS_REJECTED',
'ACCESS_GRANTED',

) NOT NULL,

-- MANAGER APPROVAL
manager_approver_id BIGINT,

-- DEVOPS APPROVAL
devops_approver_id BIGINT,

duration_hours INT DEFAULT 24,
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,

FOREIGN KEY (requester_id) REFERENCES users(id),
FOREIGN KEY (manager_approver_id) REFERENCES users(id),
FOREIGN KEY (devops_approver_id) REFERENCES users(id)
);

```

---

## **5.7 audit\_logs**

```

CREATE TABLE audit_logs (
id BIGINT PRIMARY KEY AUTO_INCREMENT,
access_request_id BIGINT NOT NULL,
event_type VARCHAR(100) NOT NULL,
event_data JSON,
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
FOREIGN KEY (access_request_id) REFERENCES access_requests(id)
);

```

## **5.8 follow up**

```

CREATE TABLE followup_questions (
id BIGINT PRIMARY KEY AUTO_INCREMENT,

```

```

access_request_id BIGINT NOT NULL,

question TEXT NOT NULL,
answer TEXT,

status ENUM(
    'ASKED',
    'ANSWERED',
    'INVALID'
) NOT NULL DEFAULT 'ASKED',

created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
answered_at TIMESTAMP,

FOREIGN KEY (access_request_id)
    REFERENCES access_requests(id)
);

```

---

## 6. LLM Design

### 6.1 LLM Input JSON

```
{
  "reason": "Need access to upload, download, and delete logs in S3",
  "allowedServices": ["S3"],
  "allowedActionGroups": {
    "S3": [
      "READ_OBJECTS",
      "UPLOAD_OBJECTS",
      "DELETE_OBJECTS",
      "LIST_BUCKET",
      "MANAGE_BUCKET"
    ]
  }
}
```

```
}
```

## 6.2 Standard Action Groups (S3)

```
{
  "S3": {
    "READ_OBJECTS": [
      "s3:GetObject",
      "s3:GetObjectVersion"
    ],
    "UPLOAD_OBJECTS": [
      "s3:PutObject",
      "s3>CreateMultipartUpload",
      "s3:UploadPart",
      "s3:CompleteMultipartUpload",
      "s3:AbortMultipartUpload"
    ],
    "DELETE_OBJECTS": [
      "s3>DeleteObject",
      "s3>DeleteObjectVersion"
    ],
    "LIST_BUCKET": [
      "s3>ListBucket"
    ],
    "MANAGE_BUCKET": [
      "s3>CreateBucket",
      "s3>DeleteBucket",
      "s3:GetBucketPolicy",
      "s3:PutBucketPolicy",
      "s3:GetBucketVersioning",
      "s3:PutBucketVersioning"
    ]
  }
}
```

### **6.3 LLM Output JSON (Follow-up Required)**

```
{  
  "needFollowup": true,  
  "followupQuestion": "Please provide the S3 bucket ARN(s) and optional object prefix.",  
  "services": ["S3"],  
  "actionGroups": [  
    "READ_OBJECTS",  
    "UPLOAD_OBJECTS",  
    "DELETE_OBJECTS",  
    "LIST_BUCKET"  
  ]  
}
```

### **6.4 LLM Output JSON (Complete)**

```
{  
  "needFollowup": false,  
  "services": ["S3"],  
  "actionGroups": [  
    "READ_OBJECTS",  
    "UPLOAD_OBJECTS",  
    "DELETE_OBJECTS",  
    "LIST_BUCKET"  
  ],  
  "resolvedResources": {  
    "S3": {  
      "bucketArn": "arn:aws:s3:::logs-bucket",  
      "objectArn": "arn:aws:s3:::logs-bucket/logs/*"  
    }  
  }  
}
```

#### **Changes made**

**call LLM**

**IF follow-up needed:**

**store question in followup\_questions**

**return followupQuestion to FE**

### Pseudo code

```
llmResponse = callLLM(reason)

if llmResponse.needFollowup == true:
    insert into followup_questions (
        access_request_id,
        question,
        status
    ) values (
        requestId,
        llmResponse.followupQuestion,
        'ASKED'
    )
```

## 7. IAM Policy & CLI Output (From Action Groups)

---

### 7.1 Generated IAM Policy JSON

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "ListLogsBucket",
            "Effect": "Allow",
            "Action": [
                "s3>ListBucket"
            ],
            "Resource": "arn:aws:s3:::logs-bucket"
        },
        {
            "Sid": "ReadWriteDeleteLogs",
            "Effect": "Allow",

```

```
"Action": [
    "s3:GetObject",
    "s3:GetObjectVersion",
    "s3:PutObject",
    "s3>CreateMultipartUpload",
    "s3:UploadPart",
    "s3:CompleteMultipartUpload",
    "s3:AbortMultipartUpload",
    "s3:DeleteObject",
    "s3:DeleteObjectVersion"
],
"Resource": "arn:aws:s3:::logs-bucket/logs/*"
}
]
}
```

---

## 7.2 AWS CLI Command (Generated)

```
aws iam create-policy \
--policy-name TempS3LogsFullAccess_001 \
--policy-document file://policy.json
```

---

## 8. Phase-1 Constraints

- No automatic AWS execution
  - No policy attachment
  - No revoke workflow
  - Manual DevOps execution only
- 

## 9. Phase-1 Completion Criteria

Phase-1 is considered complete when:

- Requests are fully structured using LLM assistance
  - Follow-up questions resolve missing data
  - Approvals are enforced
  - IAM policy & CLI are generated
  - All actions are auditable
- 

## 10. Conclusion

This design satisfies all Phase-1 PRD requirements while maintaining strong governance, deterministic LLM behavior, and extensibility for Phase-2 enhancements.

API : Changes asked -

### 1. Authentication (Common for all APIs)

- All APIs are authenticated via **SSO**
- User identity is derived from request headers (example):

X-User-Email: user@company.com

X-User-Role: EMPLOYEE | MANAGER | DEVOPS

---

### 2. Create Access Request API

#### Endpoint

POST /api/v1/access-requests

## Purpose

- Accept initial access request
  - Invoke LLM to suggest services/actions
  - take answer
  - update followup\_questions
  - continue flow
- 

## Request Body

```
{  
  "reason": "Need access to upload logs to S3",  
  "awsAccount": "ZINKA",  
  "services": ["S3"],  
  "actionGroups": [  
    "READ_OBJECTS"  
  ],  
  "resources": {  
    "bucketArn": "arn:aws:s3:::logs-bucket",  
    "objectArn": "arn:aws:s3:::logs-bucket/*"  
  }  
}
```

---

## **Response – Follow-up Required**

```
{  
  "requestId": 101,  
  "needFollowup": true,  
  "followupQuestion": "Please provide the S3 bucket ARN(s).",  
  "services": ["S3"],  
  "actionGroups": [  
    "READ_OBJECTS",  
    "UPLOAD_OBJECTS"  
  ]  
}
```

---

When `needFollowup = true`, the system persists the follow-up question in the `followup_questions` entity for tracking and audit purposes. The question is returned in the API response for frontend display.

## **Response – No Follow-up Required**

```
{  
  "requestId": 101,  
  
  "needFollowup": false,  
  
  "services": ["S3"],  
  
  "actionGroups": [  
    "READ_OBJECTS",  
    "UPLOAD_OBJECTS"  
  ],  
  
  "status": "CREATED"  
}
```

---

### **3. Submit Follow-up Answer API**

#### **Endpoint**

POST /api/v1/access-requests/{requestId}/followup/{followup\_id}

#### **Purpose**

- Accept user's response to LLM follow-up question
  - Re-evaluate request completeness
- 

## Request Body

```
{  
  "answer": {  
    "S3": {  
      "bucketArn": "arn:aws:s3:::logs-bucket",  
      "objectArn": "arn:aws:s3:::logs-bucket/*"  
    }  
  }  
}
```

---

## Response

```
{  
  "requestId": 101,  
  "status": "CREATED",  
  "services": ["S3"],  
  "actionGroups": [  
    ...  
  ]  
}
```

```
"READ_OBJECTS",
"UPLOAD_OBJECTS"
]
}
```

Backend behaviour :

1. User submits follow-up answer
2. Backend updates followup\_questions.answer
3. Backend calls LLM again with:
  - original reason
  - previous LLM context
  - structured follow-up answer
4. LLM returns final structured output
5. Backend updates:
  - access\_requests.resource\_arns
  - access\_requests.services
  - access\_requests.status = CREATED

---

#### 4. Get Access Request Details API

## **Endpoint**

GET /api/v1/access-requests/{requestId}

---

## **Purpose**

- Fetch request details for UI display
  - Used by Employee, Manager, and DevOps views
- 

## **Response**

```
{  
  "requestId": 101,  
  "requesterEmail": "user@company.com",  
  "awsAccount": "ZINKA",  
  "reason": "Need access to upload logs to S3",  
  "services": ["S3"],  
  "actionGroups": ["Upload Objects"],  
  "resourceArns": ["arn:aws:s3:::logs-bucket/*"],  
  "status": "CREATED",  
  "createdAt": "2026-01-19T10:30:00Z"  
}
```

---

## **5. Manager Approval API**

## **Endpoint**

POST /api/v1/access-requests/{requestId}/manager-approval

---

## **Purpose**

- Manager approves or rejects request

---

## **Request Body**

```
{  
  "decision": "APPROVED",  
  "reason": "Required for log processing"  
}
```

---

## **Response**

```
{  
  "requestId": 101,  
  "status": "MANAGER_APPROVED"  
}
```

---

## **6. DevOps Approval API**

### **Endpoint**

POST /api/v1/access-requests/{requestId}/devops-approval

## Purpose

- DevOps reviews and approves request
  - Triggers IAM policy + CLI generation
- 

## Request Body

```
{  
  "decision": "APPROVED",  
  "reason": "Scope looks correct"  
}
```

---

## Response

```
{  
  "requestId": 101,  
  "status": "DEVOPS_APPROVED",  
  "iamPolicy": {  
    "Version": "2012-10-17",  
    "Statement": [  
      {  
        "Effect": "Allow",
```

```
        "Action": ["s3:PutObject"],  
        "Resource": ["arn:aws:s3:::logs-bucket/*"]  
    },  
]  
,  
"awsCliCommand": "aws iam create-policy --policy-name TempS3Access_101  
--policy-document file://policy.json"  
}
```

---

## 7. Mark Access Granted API

### Endpoint

POST /api/v1/access-requests/{requestId}/mark-granted

---

### Purpose

- DevOps confirms manual execution of CLI command
- 

### Request Body

{}

---

### Response

```
{  
    "requestId": 101,  
    "status": "ACCESS_GRANTED"  
}
```

---

## 8. Error Response (Common)

```
{  
    "errorCode": "INVALID_REQUEST",  
    "message": "Missing required field: awsAccount"  
}
```

---

## Spring Boot Layered Architecture

```
com.company.awsaccess  
  └── AwsAccessApplication  
      ├── controller  
      │   ├── AccessRequestController  
      │   ├── FollowupController  
      │   └── ApprovalController  
      └── service  
          ├── AccessRequestService  
          ├── LlmAssistService  
          ├── FollowupService  
          └── ApprovalService
```

```
    └── PolicyGenerationService
    └── AuditLogService

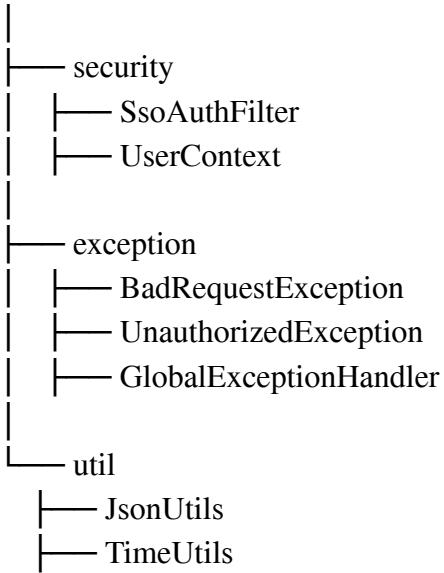
    └── repository
        └── UserRepository
        └── ServiceRepository
        └── ActionGroupRepository
        └── AwsActionRepository
        └── AccessRequestRepository
        └── ApprovalRepository
        └── FollowupQuestionRepository
        └── AuditLogRepository

    └── model (Entity)
        └── User
        └── ServiceEntity
        └── ActionGroup
        └── AwsAction
        └── ActionGroupAction
        └── AccessRequest
        └── Approval
        └── FollowupQuestion
        └── AuditLog

    └── dto
        └── request
            └── CreateAccessRequestDto
            └── FollowupAnswerDto
            └── ApprovalDecisionDto

        └── response
            └── AccessRequestResponseDto
            └── FollowupRequiredResponseDto
            └── IamPolicyResponseDto

    └── llm
        └── LlmClient
        └── LlmRequestBuilder
        └── LlmResponseParser
```



## 2. Controller Layer (API Layer)

### 2.1 AccessRequestController

Handles **request creation** and **fetching request details**

```

@RestController
@RequestMapping("/api/v1/access-requests")
public class AccessRequestController {

```

#### APIs handled:

- POST /api/v1/access-requests
- GET /api/v1/access-requests/{requestId}
- POST /api/v1/access-requests/{requestId}/mark-granted

### 2.2 FollowupController

Handles **LLM follow-up lifecycle**

```
@RestController  
@RequestMapping("/api/v1/access-requests")  
public class FollowupController {
```

**APIs handled:**

- POST /{requestId}/followup/{followupId}

This separation satisfies:

“Create a separate entity and flow for follow-up questions”

### **2.3 ApprovalController**

```
@RestController  
@RequestMapping("/api/v1/access-requests")  
public class ApprovalController {
```

**APIs handled:**

- POST /{requestId}/manager-approval
  - POST /{requestId}/devops-approval
- 

## **3. Service Layer (Business Logic)**

This is **where ALL PRD rules live.**

---

### **3.1 AccessRequestService**

#### **Responsibilities**

- Create access request

- Persist initial request
- Call LLM
- Control request state transitions

```
public class AccessRequestService {
```

### **Key methods**

```
createRequest(CreateAccessRequestDto dto)  
getRequestById(Long requestId)  
markAccessGranted(Long requestId)
```

---

## **3.2 LlmAssistService**

### **Single responsibility: LLM interaction**

```
public class LlmAssistService {
```

### **Responsibilities**

- Build LLM input JSON
- Call LLM
- Validate deterministic output
- Parse response

```
LlmResponse callLlm(LlmInput input)
```

This service is used in:

- Create Access Request
  - Submit Follow-up Answer
- 

### **3.3 FollowupService**

```
public class FollowupService {
```

#### **Responsibilities**

- Persist follow-up question
- Persist follow-up answer
- Trigger LLM re-evaluation

```
createFollowupQuestion(requestId, question)  
answerFollowup(followupId, structuredAnswer)
```

Directly maps to:

```
followup_questions table  
status = ASKED → ANSWERED
```

---

### **3.4 ApprovalService**

```
public class ApprovalService {
```

#### **Responsibilities**

- Validate approver role
- Record approval decision

- Move request state

```
managerApproval(requestId, decision)  
devopsApproval(requestId, decision)
```

---

### **3.5 PolicyGenerationService**

```
public class PolicyGenerationService {
```

#### **Responsibilities**

- Convert action groups → IAM actions
- Generate IAM policy JSON
- Generate AWS CLI command (text only)

```
generatePolicy(accessRequest)  
generateCliCommand(policy)
```

No AWS execution (Phase-1 constraint enforced here).

---

### **3.6 AuditLogService**

```
public class AuditLogService {
```

#### **Responsibilities**

- Write immutable audit logs for:
  - request creation

- LLM decisions
  - follow-up asked
  - follow-up answered
  - approvals
  - policy generation
  - access granted
- 

## 4. Repository Layer (Persistence)

Each table has **one repository**, no shortcuts.

public interface AccessRequestRepository extends JpaRepository<AccessRequest, Long>

Same pattern for:

- UserRepository
- ApprovalRepository
- FollowupQuestionRepository
- AuditLogRepository
- ServiceRepository
- ActionGroupRepository
- AwsActionRepository

---

## 5. Entity Layer (Exact DB Mapping)

Entities map **1-to-1** with your DDL.

Examples:

### **AccessRequest Entity**

```
@Entity  
@Table(name = "access_requests")  
public class AccessRequest {
```

Fields:

- id
- requesterId
- awsAccount
- reason
- services (JSON)
- resourceArns (JSON)
- status
- durationHours
- createdAt

---

### **FollowupQuestion Entity**

```
@Entity
```

```
@Table(name = "followup_questions")
public class FollowupQuestion {
```

Fields:

- id
  - accessRequestId
  - question
  - answer
  - status
  - createdAt
  - answeredAt
- 

## 6. DTO Layer (Frontend Contract)

This is **critical for FE sharing**.

### CreateAccessRequestDto

```
reason
awsAccount
services[]
actionGroups[]
resources{}
```

### FollowupAnswerDto

```
answer {
  S3 {
    bucketArn
```

```
    objectArn  
}  
}
```

## **AccessRequestResponseDto**

```
requestId  
status  
services  
actionGroups  
followupQuestion
```

---

## **7. Security Layer (SSO)**

```
public class SsoAuthFilter extends OncePerRequestFilter
```

Reads headers:

X-User-Email  
X-User-Role

Populates:

```
UserContext.set(email, role)
```

Used across services for:

- authorization
  - approvals
  - audit logs
-

## 8. Exception Handling

```
@ControllerAdvice  
public class GlobalExceptionHandler {
```

Handles:

- invalid request
- missing follow-up
- unauthorized approval
- invalid state transitions

Maps to:

```
{  
  "errorCode": "INVALID_REQUEST",  
  "message": "..."  
}
```

---